Project 1

Jane Kim*

Physics 480: Computational Physics
(Dated: March 5, 2018)

This project aims to approximate the solution to the 1D Poisson equation with Dirichlet boundary conditions, $-u''(x) = f(x)$, by rewriting the problem as a set of linear equations, $A\vec{v} = \vec{b}$. Four algorithms, developed with different initial assumptions about the form of A, were implemented in C++ and were tested against a known analytic solution. The methods were also timed and the relative errors between the approximate and exact solutions were studied. The method using standard LU decomposition was the slowest among the four. The other three had comparable computation times, but the simplest algorithm proved to be the most stable for large n.

I. INTRODUCTION

The goal of this project was to develop and compare several methods of solving the one-dimensional Poisson equation with Dirichlet boundary conditions. More explicitly, we aimed to solve

$$-u''(x) = f(x), \quad x \in (0,1), \quad u(0) = u(1) = 0, \quad (1)$$

for some given function $f(x)$. We first rewrote this problem as a set of linear equations by approximating the second derivative of $u(x)$ using the discretized approximation to $u(x)$, namely $v_i = u(x_i)$, where $x_i = ih$ is defined for $i = 0, ..., n+1$ and $h = 1/(n+1)$ is the step size. The Poisson equation then becomes

$$-\frac{v_{i-1} - 2v_i + v_{i+1}}{h^2} = f(x_i), \quad i = 1, ..., n. \quad (2)$$

Hence we obtained a set of n linear equations given by

$$A\vec{v} = \vec{b},$$

$$\begin{bmatrix} 2 & -1 & & & \\ -1 & 2 & -1 & & \\ & \ddots & \ddots & \ddots & \\ & & -1 & 2 & -1 \\ & & & -1 & 2 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_{n-1} \\ v_n \end{bmatrix} = h^2 \begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_{n-1}) \\ f(x_n) \end{bmatrix}. \quad (3)$$

There exists many different methods for solving Equation (3) for $\vec{v}$, but the efficiency of the methods discussed in this report were mainly dependent on the initial assumptions of the form of A. Four different algorithms were implemented according to the following assumptions: (i) A is dense, (ii) A is tridiagonal, (iii) A is symmetric and tridiagonal, and (iv) A has the form shown in Equation (3).

Section II describes the four different algorithms. In Section III, the tests and timings of the algorithms are discussed. Results are presented in Section IV, along with some commentary. And finally, Section V concludes this report.

II. METHODS

A. LU Decomposition with Armadillo

This method uses standard LU decomposition to solve Equation (3) with no prior assumptions about the form of A. While this algorithm is simple to write (especially with Armadillo) and is more efficient than Gaussian elimination, the fact that A has more zeros than non-zero elements results in a relatively long computation time and storage issues for large n. For instance, running the program with $n = 10^5$ resulted in an early termination due to lack of available memory. The number of floating point operations (FLOPS) for LU decomposition is $\sim \frac{2}{3}n^3$ for large n.

B. General Tridiagonal Matrix

By assuming A is tridiagonal, we can store the three main bands as $n-$dimensional vectors instead of storing the entire matrix. The main diagonal vector was denoted as $\vec{d}$, and the lower and upper diagonals were denoted as $\vec{c}$ and $\vec{e}$, respectively. The algorithm was a simplified form of Gaussian elimination and consisted of two steps. In the forward substitution step, the lower diagonal elements c_i were eliminated for $i = 2, ..., n$ by Gaussian elimination, leaving 1's along the diagonals. With this construction, only $\vec{e}$ and $\vec{b}$ needed to be changed:

$$e_i = \begin{cases} \frac{e_i}{d_i}, & i = 1, \\ \frac{e_i}{d_i - e_{i-1}c_i}, & i = 2, ..., n, \end{cases} \quad (4)$$

$$b_i = \begin{cases} \frac{b_i}{d_i}, & i = 1, \\ \frac{b_i - b_{i-1}c_i}{d_i - e_{i-1}c_i}, & i = 2, ..., n. \end{cases} \quad (5)$$

Notice that the denominators for the $i > 1$ case is the same for both e_i and b_i. So this value was precalculated to reduce the number of floating point operations. The forward substitution step, therefore, consists of $\sim 6n$ FLOPS for large n.

The backward substitution step solves for $\vec{v}$ using the new values of $\vec{e}$ and $\vec{b}$:

$$v_i = \begin{cases} b_i, & i = n, \\ b_i - e_i v_{i+1}, & i = n-1, ..., 1, \end{cases} \quad (6)$$

thus it uses $\sim 2n$ FLOPS. So the algorithm employs $\sim 8n$ FLOPS all together.

C. Symmetric Tridiagonal Matrix

We can simplify the previous algorithm even further. If we assume that $\vec{c} = \vec{e}$, then the lower diagonal elements can be eliminated while keeping the upper diagonal elements fixed. Then the forward substitution step is given by

$$d_i = \begin{cases} d_i, & i = 1, \\ d_i - \frac{c_{i-1}^2}{d_{i-1}}, & i = 2, ..., n, \end{cases} \quad (7)$$

$$b_i = \begin{cases} b_i, & i = 1, \\ b_i - \frac{c_{i-1} b_{i-1}}{d_{i-1}}, & i = 2, ..., n. \end{cases} \quad (8)$$

Again, the factor $\frac{c_{i-1}}{d_{i-1}}$ can be precalculated so that the forward substitution step uses $\sim 5n$ floating point operations.

The backward substitution step is as follows:

$$v_i = \begin{cases} \frac{b_i}{d_i}, & i = n, \\ \frac{b_i - c_i v_{i+1}}{d_i}, & i = n-1, ..., 1. \end{cases} \quad (9)$$

This step requires $\sim 3n$ FLOPS so that the overall algorithm uses about the same number of FLOPS as the general tridiagonal case.

D. Simple Tridiagonal Matrix

Finally, let $d_i = 2$ and $c_i = -1$ for $i = 1, ..., n$ in the previous algorithm. By inspection, we find that the forward substitution step for $\vec{d}$ follows a predictable pattern, where $d_i = \frac{i+1}{i}$. So we can precalculate the entire vector $\vec{d}$ in the set-up step, rather than the forward substitution step. Then the forward and backward substitution steps only consist of

$$b_i = \begin{cases} b_i, & i = 1, \\ b_i + \frac{b_{i-1}}{d_{i-1}}, & i = 2, ..., n, \end{cases} \quad (10)$$

$$v_i = \begin{cases} \frac{b_i}{d_i}, & i = n, \\ \frac{b_i + v_{i+1}}{d_i}, & i = n-1, ..., 1, \end{cases} \quad (11)$$

which together uses $\sim 4n$ FLOPS. Clearly, this is the least number of floating point operations among the four methods discussed in this report, but this simple algorithm is the least versatile, as it solves Equation (3) only.

Note: Since many programming languages use zero-based numbering, one may need to alter the equation for d_i slightly: $d_i = \frac{i+2}{i+1}$, $i = 0, ..., n-1$.

III. IMPLEMENTATION

All four programs require a maximum power of 10 and a file name in the command line. Then the programs loop through successive powers of 10 for the number of grid points n. In each pass of the loop, a final file name is designated with the corresponding power of 10.

Each program was tested with $f(x) = 100e^{-10x}$ so that Equation (1) had a known analytic solution $u(x) = 1 - (1 - e^{-10})x - e^{-10x}$ with which we could compare the discretized approximation v_i. It is important to note that the set up step was not included in the timed portion of the program. Only the algorithm portion, either forward and backward substitution or call to function, of each program was timed.

In addition, a vector $\vec{err}$ was created to store the relative errors between $u(x_i)$ and v_i. Armadillo allows for easy extraction of the maximum, minimum, and average relative error values. The three programs in which Armadillo was not used included explicit calculations of these values. Finally, the function inputs x_i, the approximate solution v_i, the exact solution $u(x_i)$, and the logarithm of the relative errors were printed to file. The number of grid points n, the time used for computation, and the maximum, minimum, and average relative error values were also printed to the beginning of each file.

For methods B-D, example file outputs for $n = 10^1, ..., 10^6$ can be found in the 'benchmark' folder. For method A, $n > 10^3$ proved to be slow and problematic, so only three benchmark calculations were provided.

IV. RESULTS AND ANALYSIS

The computation times for all four programs are shown in Table 1. LU decomposition (A) requires the most computation time and is unable to produce the solutions for $n > 10^3$ in a relatively timely manner. The times for the general tridiagonal matrix (B) and the symmetric tridiagonal matrix (C) algorithms are comparable to each other, but a few orders of magnitude faster than LU decomposition. The simple tridiagonal matrix algorithm (D) produces the fastest results out of the four methods,

but only by a small margin considering methods B and C require twice as many FLOPS as method D.

n	10	10^2	10^3	10^4	10^5	10^6
A	$2.1E-4$	$1.4E-4$	0.288	N/A	N/A	N/A
B	$2.0E-6$	$6.0E-6$	$4.4E-5$	$4.1E-4$	$2.4E-3$	0.023
C	$2.0E-6$	$4.0E-6$	$3.4E-5$	$3.6E-4$	$3.1E-3$	0.030
D	$1.0E-6$	$3.0E-6$	$3.1E-5$	$3.2E-4$	$2.5E-3$	0.024

Table 1. Time used for computation (in seconds) for methods A, B, C, and D and for $n = 10 - 10^6$. Benchmark calculations can be found in the 'benchmark' folder under 'report'.

The relative errors between the discretized approximation and the closed-form solution are presented in Table 2. For $n \leq 10^4$, the logarithms of the errors are largely the same and behave as expected for this particular approximation of the second derivative, i.e. $\log(err)$ vs. h has a slope of 2. However, the table suggests methods B and C begin to succumb to rounding errors for $n > 10^4$ since the slope just starts to drop. Then for $n = 10^6$ the errors are much larger. Thus, methods B and C are optimal for $n \approx 10^4 - 10^5$. On the other hand, method D continues to have a slope of 2 for all n values tested, so the optimal range for n could be even greater than 10^6.

n	10	10^2	10^3	10^4	10^5	10^6
A	-1.180	-3.088	-5.080	N/A	N/A	N/A
B	" "	" "	" "	-7.079	-8.889	-6.167
C	" "	" "	" "	" "	" "	" "
D	" "	" "	" "	" "	-9.079	-11.083

Table 2. The logarithm of the average relative error for methods A, B, C, and D and for $n = 10 - 10^6$. The quotations indicate repeated values from above.

The discretized approximation appears to converge quite quickly to the analytic solution as n increases, considering the two look almost identical for $n > 100$ (See Figure 1). Figure 1 was specifically generated using the output from the simple tridiagonal matrix algorithm (D), but analogous plots from the other three algorithms are numerically identical for $n = 10, 100, 1000$.

V.   CONCLUSION

Of the four methods studied here, the simple tridiagonal matrix algorithm (D) was the fastest and most stable for large n. However, it was limited in scope because it was developed for this particular discretization of the second derivative. If we had instead included additional terms in Equation (2), the matrix A would no longer be tridiagonal and only LU decomposition would have the versitility to solve Equation (3).

This study could be further expanded to include iterative methods, such as the Gauss-Seidel method discussed in class. It may also be interesting to test other functions $f(x)$ that have analytic solutions.

I tried to get my references to show up, but the LaTeXpackage I've used in the past is no longer working for some reason... I would mainly like to cite the class lecture notes.
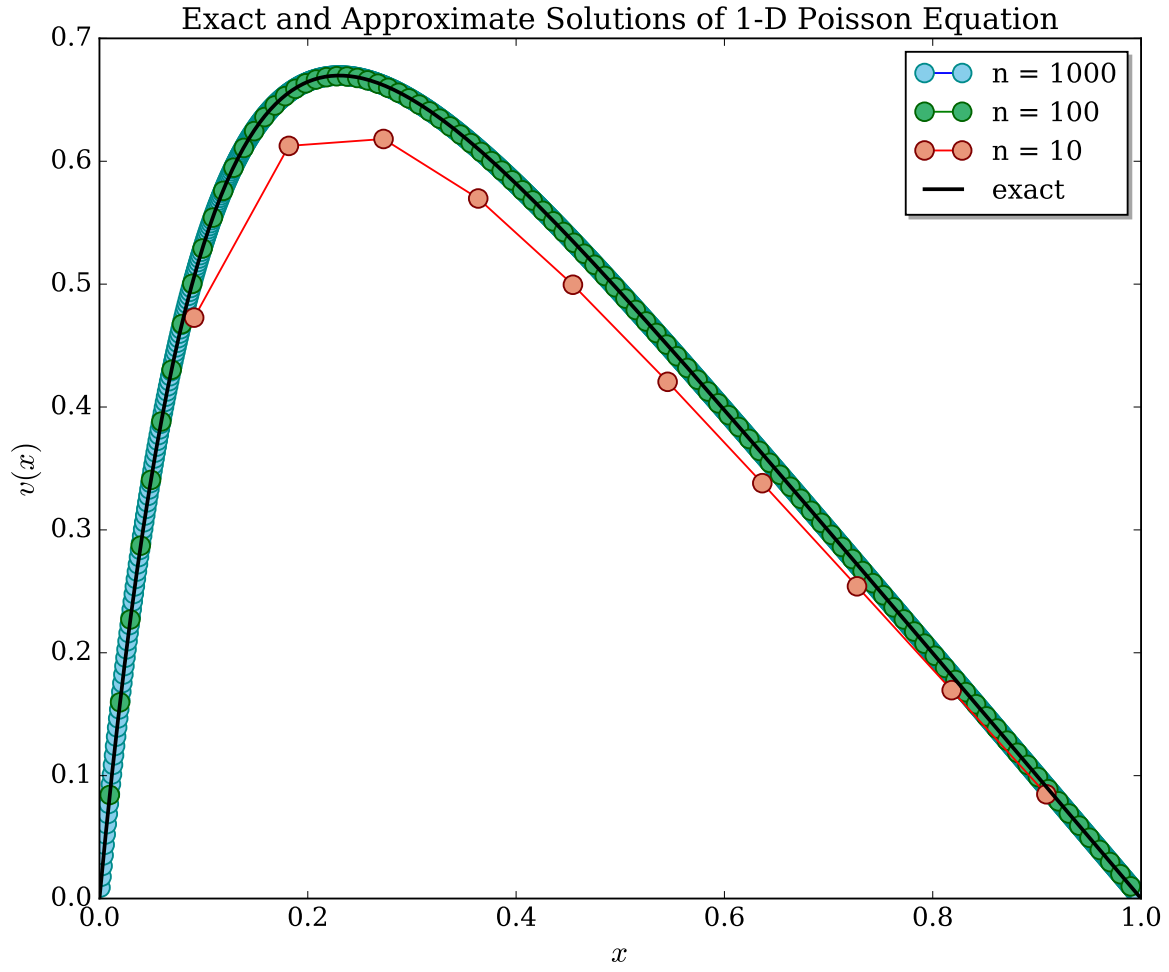
Figure 1. The discretized solution v_i for various n are plotted alongside the closed-form solution $u(x)$. For $n > 100$, the approximate and exact solutions are almost indistinguishable.